

Algorithms (I)

Introduction

Guoqiang Li

School of Software, Shanghai Jiao Tong University

Instructor and Teaching Assistants

- Guoqiang LI
 - Homepage: <http://basics.sjtu.edu.cn/~liguoqiang>
 - Course page:
<http://basics.sjtu.edu.cn/~liguoqiang/teaching/Galgo17/index.htm>
 - Email: li.g@outlook.com
 - Office: Rm. 1212, Building of Software
 - Phone: [3420-4167](tel:3420-4167)
- TA:
 - **Qizhe YANG**: yangqizhe1994 (AT) gmail (DOT) com
 - **Yuwei WANG**: wangyuwei95 (AT) qq (DOT) com
- Office hour: [Tue. 14:00-17:00 @ SEIEE 3-327 & Software Building 3203](#)

Notification

- Students who take this lecture are assumed to have a solid background of algorithms.
- **Principle of Algorithms.**
- Students are **NOT** expected to give a presentation in this lecture.

Algorithm Design

- Basic methodologies:
 - Algorithms on Lists, Trees and Graphs
 - Divide and Conquer
 - Master Theorem
 - Recursion
- Advanced topics:
 - Dynamic Programming
 - Greedy Algorithm
 - Linear Programming
 - Approximation Algorithm
 - Randomized Algorithm
 - Computational Geometry
 - ...

Algorithm Analysis

- Big-*O* Notation
- Advanced Methodology:
 - Probability Analysis
 - Amortized Analysis
 - Competition Analysis

Standard Algorithms


- Sorting
- Searching & Hashing
- Strongly connected components
- Finding shortest paths in graphs
- Minimum spanning trees in graphs
- Matchings in bipartite graphs
- Maximum flows in networks

Data Structure

- Link lists
- Trees, graphs
- Kripke structure, automata
- Priority queue
- Disjoint set

Reference Book


- **Algorithms**
 - Sanjoy Dasgupta
 - San Diego Christos Papadimitriou
 - Umesh Vazirani
 - McGraw-Hill, 2007.



Reference Book


- **Introduction to Algorithms**

- Thomas H. Cormen
- Charles E. Leiserson
- Ronald L. Rivest
- Clifford Stein
- The MIT Press (3rd edition), 2009.




Reference Book

- **Approximation Algorithms**
 - Vijay V. Vazirani
 - Springer-Verlag, 2004



Reference Book

- **The Design of Approximation Algorithms**
 - David P. Williamson
 - David B. Shmoys
 - Cambridge University Press, 2011.



Scoring Policy

- 10% Attendance.
- 20% Homework.
 - Four assignments.
 - Each one is 5pts.
 - Work out individually.
 - Each assignment will be evaluated by *A, B, C, D, F* (Excellent(5), Good(5), Fair(4), Delay(3), Fail(0))
- 70% Final exam.

Any Questions?

Typical Algorithms

Sorting

Sorting

Input: A sequence of n numbers $\langle a_1, a_2, \dots, a_n \rangle$.

Output: A permutation (reordering) $\langle a'_1, a'_2, \dots, a'_n \rangle$ of the input sequence such that $a'_1 \leq a'_2 \leq \dots \leq a'_n$

Various Sorts

- Insert Sort
- Bubble Sort
- Heap Sort
- Quick Sort
- Merge Sort

Merge Sort

The Algorithm

MERGESORT ($a[1 \dots n]$)

An array of numbers $a[1 \dots n]$;

if $n > 1$ **then**

 return (MERGE (MERGESORT ($a[1 \dots \lfloor n/2 \rfloor]$),

 MERGESORT ($a[\lfloor n/2 \rfloor + 1 \dots n]$)));

else return (a);

end

MERGE ($x[1 \dots k], y[1 \dots l]$)

if $k = 0$ **then** return $y[1 \dots l]$;

if $l = 0$ **then** return $x[1 \dots k]$;

if $x[1] \leq y[1]$ **then**

 return ($x[1] \circ \text{MERGE} (x[2 \dots k], y[1 \dots l])$);

else return ($y[1] \circ \text{MERGE} (x[1 \dots k], y[2 \dots l])$);

end

An Iterative Version

```
ITERATIVE-MERGESORT ( $a[1 \dots n]$ )  
An array of numbers  $a[1 \dots n]$ ;  
 $Q = [ ]$  empty queue;  
for  $i = 1$  to  $n$  do  
    | Inject ( $Q, [a]$ );  
end  
while  $|Q| > 1$  do  
    | Inject ( $Q, \text{MERGE}(\text{Eject}(Q), \text{Eject}(Q))$ );  
end  
return ( $\text{Eject}(Q)$ );
```

The Time Analysis

- The recurrence relation:

$$T(n) = 2T(n/2) + O(n)$$

- By Master Theorem:

$$T(n) = O(n \log n)$$

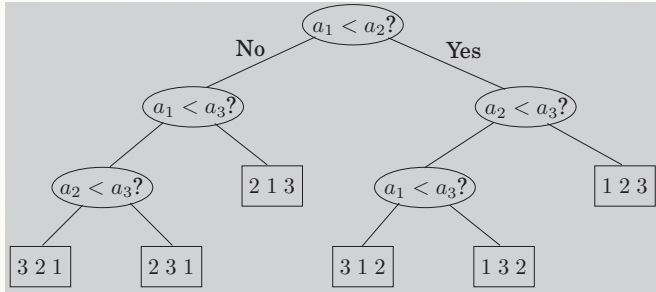
Master Theorem

If $T(n) = aT(\lceil n/b \rceil) + O(n^d)$ for some constants $a > 0$, $b > 1$ and $d \geq 0$, then

$$T(n) = \begin{cases} O(n^d) & \text{if } d > \log_b a \\ O(n^d \log n) & \text{if } d = \log_b a \\ O(n^{\log_b a}) & \text{if } d < \log_b a \end{cases}$$

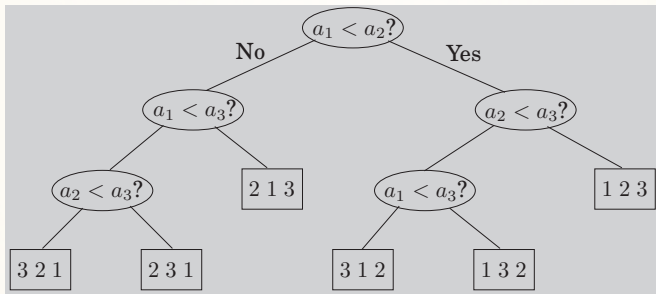
Can we do better?

Sorting



- A sorting algorithm can be depicted as a **decision tree**.
- The **depth** of the tree - the number of comparisons on the longest path from root to leaf, is the worst-case time complexity of the algorithm.
- Assume n elements. Each of its leaves is labeled by a **permutation** of $\{1, 2, \dots, n\}$.

Sorting



- Every permutation must appear as the label of a leaf.
- This is a binary tree with $n!$ leaves.
- So ,the depth of the tree - and the complexity of the algorithm - must be at least


$$\log(n!) \approx \log(\sqrt{\pi(2n + 1/3)} \cdot n^n \cdot e^{-n}) = \Omega(n \log n)$$

Shortest Path in Graph

Lengths on Edges

- **BFS** treats all edges as having the same length.
- It is rarely true in applications where shortest paths are to be found.
- Every edge $e \in E$ with a length l_e .
- If $e = (u, v)$, we will sometimes also write

$$l(u, v) \quad \text{or} \quad l_{uv}$$



Dijkstra's Algorithm

An Adaption of Breadth-First Search



- **BFS** finds shortest paths in any graph whose edges have unit length.
- **Q:** Can we adapt it to a more general graph $G = (V, E)$ whose edge lengths l_e are positive integers?
- **A simple trick:** For any edge $e = (u, v)$ of E , replace it by l_e edges of length 1, by adding $l_e - 1$ dummy nodes between u and v . It might take time

$$O(|V| + \sum_{e \in E} l_e)$$


- It is **bad** in case we have edges with high length.

Alarm Clocks


- Set an alarm clock for node s at time 0 .
- Repeat until there are no more alarms:
- Say the next alarm goes off at time T , for node u . Then:
 - The distance from s to u is T .
 - For each neighbor v of u in G :
 - If there is no alarm yet for v , set one for time $T + l(u, v)$.
 - If v 's alarm is set for later than $T + l(u, v)$, then reset it to this earlier time.




An Example




A: 0	D: ∞
B: 4	E: ∞
C: 2	




A: 0	D: 6
B: 3	E: 7
C: 2	



A: 0	D: 5
B: 3	E: 6
C: 2	



A: 0	D: 5
B: 3	E: 6
C: 2	



Priority Queue

- **Priority queue** is a data structure usually implemented by heap.
 - **Insert:** Add a new element to the set.
 - **Decrease-key:** Accommodate the decrease in key value of a particular element.
 - **Delete-min:** Return the element with the smallest key, and remove it from the set.
 - **Make-queue:** Build a priority queue out of the given elements, with the given key values. (In many implementations, this is significantly faster than inserting the elements one by one.)
- The first two let us **set alarms**, and the third tells us which alarm is **next** to go off.

Dijkstra's Shortest-Path Algorithm

DIJKSTRA (G, l, s)

input : Graph $G = (V, E)$, directed or undirected; positive edge length $\{l_e \mid e \in E\}$;
Vertex $s \in V$

output: For all vertices u reachable from s , $dist(u)$ is the set to the distance from s to u

for all $u \in V$ **do**

$dist(u) = \infty$;
 $prev(u) = nil$;

end

$dist(s) = 0$;

$H \leftarrow makequeue(V) \setminus \setminus$ using $dist$ -values as keys;

while H is not empty **do**

$u \leftarrow deletemin(H)$;

for all edge $(u, v) \in E$ **do**

if $dist(v) > dist(u) + l(u, v)$ **then**

$dist(v) = dist(u) + l(u, v)$; $prev(v) = u$;
 decreasekey(H, v);

end

end

end

Running Time

- Since `makequeue` takes at most as long as $|V|$ insert operations, we get a total of $|V|$ `deletemin` and $|V| + |E|$ `insert/decreasekey` operations.
- The time needed for these varies by implementation; for instance, a **binary heap** gives an overall running time of

$$O((|V| + |E|) \log |V|)$$

Which Heap is Best

Implementation	deletemin	insert/decreasekey	$ V \times \text{deletemin} + (V + E) \times \text{insert}$
Array	$O(V)$	$O(1)$	$O(V ^2)$
Binary heap	$O(\log V)$	$O(\log V)$	$O((V + E) \log V)$
d-ary heap	$O(\frac{d \log V }{\log d})$	$O(\frac{\log V }{\log d})$	$O(\frac{(d V + E) \log V }{\log d})$
Fibonacci heap	$O(\log V)$	$O(1)$ (amortized)	$O(V \log V + E)$

Which heap is Best

- A naive array implementation gives a respectable time complexity of $O(|V|^2)$, whereas with a binary heap we get $O((|V| + |E|) \log |V|)$. Which is preferable?
- This depends on whether the graph is **sparse** or **dense**.
 - $|E|$ is less than $|V|^2$. If it is $\Omega(|V|^2)$, then clearly the array implementation is the faster.
 - On the other hand, the binary heap becomes preferable as soon as $|E|$ dips below $|V|^2 / \log |V|$.
 - The d-ary heap is a generalization of the binary heap and leads to a running time that is a function of d . The optimal choice is $d \approx |E|/|V|$;

Edit Distance

The problem

- When a **spell checker** encounters a possible misspelling, it looks in its dictionary for other words that are close by.
 - **Q:** What is the appropriate notion of closeness in this case?
- A natural measure of the distance between two strings is the extent to which they can be **aligned**, or **matched up**.
- Technically, an alignment is simply a way of writing the strings one above the other.

S	—	N	O	W	Y
S	U	N	N	—	Y

Cost: 3

—	S	N	O	W	—	Y
S	U	N	—	—	N	Y

Cost: 5

The problem

- The **cost** of an alignment is the number of columns in which the letters differ.
- And the **edit distance** between two strings is the cost of their best possible alignment.
- **Edit distance** is so named because it can also be thought of as the **minimum number of edits**
 - **insertions**, **deletions**, and **substitutions** of characters needed to transform the first string into the second.

S	—	N	O	W	Y
S	U	N	N	—	Y

Cost: 3

—	S	N	O	W	—	Y
S	U	N	—	—	N	Y

Cost: 5

A Dynamic Programming Solution

- When solving a problem by dynamic programming, the most crucial question is, **What are the subproblems?**
- Our goal is to find the **edit distance** between two strings
 $x[1, \dots, m]$ and $y[1, \dots, n]$
- For every i, j with $1 \leq i \leq m$ and $1 \leq j \leq n$, let
 - $E(i, j)$: the **edit distance** between some prefix of the first string, $x[1, \dots, i]$, and some **prefix** of the second, $y[1, \dots, j]$.
- $E(i, j) = \min\{1 + E(i - 1, j), 1 + E(i, j - 1), \text{diff}(i, j) + E(i - 1, j - 1)\}$,
where $\text{diff}(i, j)$ is defined to be 0 if $x[i] = y[j]$ and 1 otherwise.

An Example

Edit distance between **EXPONENTIAL** and **POLYNOMIAL**, subproblem $E(4, 3)$ corresponds to the prefixes **EXPO** and **POL**. The rightmost column of their best alignment must be one of the following:

$\begin{matrix} O \\ - \end{matrix}$

 or


 $\begin{matrix} - \\ L \end{matrix}$

 or

 $\begin{matrix} O \\ L \end{matrix}$

Thus, $E(4, 3) = \min\{1 + E(3, 3), 1 + E(4, 2); 1 + E(3, 2)\}$.

(a)



(b)

		P	O	L	Y	N	O	M	I	A	L
EXPONENTIAL	0	1	2	3	4	5	6	7	8	9	10
	1	1	2	3	4	5	6	7	8	9	10
	2	2	2	3	4	5	6	7	8	9	10
	3	2	3	3	4	5	6	7	8	9	10
	4	3	2	3	4	5	5	6	7	8	9
	5	4	3	3	4	4	5	6	7	8	9
	6	5	4	4	4	4	5	6	7	8	9
	7	6	5	5	5	4	5	6	7	8	9
	8	7	6	6	6	5	5	6	7	8	9
	9	8	7	7	7	6	6	6	7	8	9
	10	9	8	8	8	7	7	7	7	6	7
11	10	9	8	9	8	8	8	8	7	6	

The Algorithm

```
for  $i = 0$  to  $m$  do  
  |  $E(i, 0) = i$ ;  
end  
for  $j = 1$  to  $n$  do  
  |  $E(0, j) = j$ ;  
end  
for  $i = 1$  to  $m$  do  
  | for  $j = 1$  to  $n$  do  
    |  $E(i, j) =$   
    |    $\min\{1 + E(i - 1, j), 1 + E(i, j - 1), \text{diff}(i, j) + E(i - 1, j - 1)\}$ ;  
    end  
  end  
end  
return ( $E(m, n)$ ) ;
```

The over running time is $O(m \cdot n)$.

Then, What is An Algorithm?

What Is An Algorithm

An **algorithm** is a procedure that consists of

- a **finite set of instructions** which,
- given an **input** from some set of possible inputs,
- enables us to obtain an **output** through a systematic execution of the instructions
- that **terminates** in a finite number of steps.

Not Enough, Maybe

What Is An Algorithm

- In these problems we are **searching** for a solution (path, tree, matching, etc.) from among an **exponential** population of possibilities.
- All these problems could in principle be solved in **exponential time** by checking through all candidate solutions, one by one.
- The quest for **algorithms** is about finding clever ways to bypass this process of **exhaustive search**, using clues from the input in order to dramatically narrow down the search space.

Is That **ALL**?

Lecture Agenda

- NP Problem
- Coping with NP Completeness
- Linear Programming
- Approximation Algorithms

Referred Materials

- [DPV07] *Algorithms*
- [CLRS09] *Introduction to Algorithms*
- [Vaz04] *Approximation Algorithms*
- [WS11] *The Design of Approximation Algorithms*
- Content of this lecture comes from section 2.3, 4.4, 4.5 and 6.3 in [DPV07].
- Suggest to read Chapter 15 of [CLRS09] and Chapter 6 in [DPV07].