

Algorithms (VII)

The Network Flow

Guoqiang Li

School of Software, Shanghai Jiao Tong University

LP and Its Dual

Primal LP

$$\begin{aligned} \max \quad & c^T \mathbf{x} \\ \text{subject to} \quad & A\mathbf{x} \leq b \\ & \mathbf{x} \geq 0 \end{aligned}$$

Dual LP

$$\begin{aligned} \min \quad & \mathbf{y}^T b \\ \text{subject to} \quad & \mathbf{y}^T A \geq c^T \\ & \mathbf{y} \geq 0 \end{aligned}$$

Primal LP:

$$\begin{aligned} \max \quad & c_1 x_1 + \cdots + c_n x_n \\ \text{subject to} \quad & a_{i1} x_1 + \cdots + a_{in} x_n \leq b_i \quad \text{for } i \in I \\ & a_{i1} x_1 + \cdots + a_{in} x_n = b_i \quad \text{for } i \in E \\ & x_j \geq 0 \quad \text{for } j \in N \end{aligned}$$

Dual LP:

$$\begin{aligned} \min \quad & b_1 y_1 + \cdots + b_m y_m \\ \text{subject to} \quad & a_{1j} y_1 + \cdots + a_{mj} y_m \geq c_j \quad \text{for } j \in N \\ & a_{1j} y_1 + \cdots + a_{mj} y_m = c_j \quad \text{for } j \notin N \\ & y_i \geq 0 \quad \text{for } i \in I \end{aligned}$$

LP and Its Dual

$$\begin{aligned}\max & x_1 + 6x_2 \\ & x_1 \leq 200 \\ & x_2 \leq 300 \\ & x_1 + x_2 \leq 400 \\ & x_1, x_2, x_3 \geq 0\end{aligned}$$

$$\begin{aligned}\min & 200y_1 + 300y_2 + 400y_3 \\ & y_1 + y_3 \geq 1 \\ & y_2 + y_3 \geq 6 \\ & y_1, y_2, y_3 \geq 0\end{aligned}$$

Shortest Path in LP

In the **shortest path problem**, we are given a weighted, directed graph $G = (V, E)$, with weight function $w : E \rightarrow \mathbb{R}$ mapping edges to real-valued weights, a source vertex s , and destination vertex t . We wish to compute the weight of a shortest path from s to t .

Shortest Path in LP

$$\min \sum_{e \in E} w_e x_e$$

$$\sum_{e \in \delta(S)} x_e \geq 1 \quad S \in \mathcal{S}$$

$$x_e \geq 0 \quad e \in E$$

$$\max \sum_{S \in \mathcal{S}} y_S$$

$$\sum_{S \in \mathcal{S}, e \in \delta(S)} y_S \leq w_e \quad e \in E$$

$$y_S \geq 0 \quad S \in \mathcal{S}$$

Complementary Slackness

Theorem

Assume LP problem (P) has a solution x^* and its dual problem (D) has a solution y^* .

- ① If $x_j^* > 0$, then the j -th constraint in (D) is binding.
- ② If the j -th constraint in (D) is not binding, then $x_j^* = 0$.
- ③ If $y_i^* > 0$, then the i -th constraint in (P) is binding.
- ④ If the i -th constraint in (P) is not binding, then $y_i^* = 0$.

The Simplex Algorithm

On each iteration, **simplex** has two tasks:

- ① **Check** whether the current vertex is **optimal** (and if so, **halt**).
- ② **Determine** where to move **next**.

As we will see, both tasks are easy if the vertex happens to be at the **origin**. And if the vertex is elsewhere, we will transform the **coordinate system** to move it to the **origin**!


Flows in Networks

Shipping Oil


- We have a network of pipelines along which oil can be sent.
- The **goal** is to ship as much oil as possible from the **source** s and the **sink** t .
- Each pipeline has a **maximum capacity** it can handle, and there are no opportunities for storing oil en route.

A Flow Example

(a)



(b)



Maximizing Flow

- The networks consist of a directed graph $G = (V, E)$; two special nodes $s, t \in V$, which are, respectively, a **source** and **sink** of G ; and **capacities** $c_e > 0$ on the edges.
- We would like to send as much oil as possible from s to t without exceeding the capacities of any of the edges.

Maximizing Flow

- A particular shipping scheme is called a **flow** and consists of a **variable** f_e for each **edge** e of the network, satisfying the following two properties:
 - ① It doesn't violate edge capacities: $0 \leq f_e \leq c_e$ for all $e \in E$.
 - ② For all nodes u except s and t , the amount of flow entering u **equals** the amount leaving

$$\sum_{(w,v) \in E} f_{wu} = \sum_{(u,z) \in E} f_{uz}$$

In other words, flow is conserved.


Maximizing Flow

- The size of a flow is the total quantity sent from s to t and, by the **conservation principle**, is equal to the quantity leaving s :

$$\text{size}(f) = \sum_{(s,u) \in E} f_{su}$$

- In short, our **goal** is to assign values to $\{f_e | e \in E\}$ that will satisfy a set of linear constraints and maximize a **linear objective function**.
- But this is a **linear program**! The maximum-flow problem reduces to linear programming.

The Example



LP

- 11 variables, one per edge.
- maximize $f_{sa} + f_{sb} + f_{sc}$
- 27 constraints:
 - 11 for nonnegativity (such as $f_{sa} \geq 0$)
 - 11 for capacity (such as $f_{sa} \leq 3$)
 - 5 for flow conservation (one for each node of the graph other than s and t , such as $f_{sc} + f_{dc} = f_{ce}$).

Another Representation

- First, introduce a **fictitious edge** of infinite capacity from t to s thus converting the flow to a circulation;
- The **objective** now is to **maximize** the flow on this edge, denoted by f_{ts} .
- The advantage of making this modification is that we can now require **flow conservation** at s and t as well.

Another Representation

$$\max f_{ts}$$

$$f_{ij} \leq c_{ij} \quad (i,j) \in E$$

$$\sum_{(w,i) \in E} f_{wi} - \sum_{(i,z) \in E} f_{iz} \leq 0 \quad i \in V$$


$$f_{ij} \geq 0 \quad (i,j) \in E$$

A Closer Look at the Algorithm


- All we know so far of the simplex algorithm is the vague **geometric intuition** that it keeps making **local moves** on the **surface of a convex** feasible region, successively improving the **objective function** until it finally reaches the **optimal solution**.
- The behavior of simplex has an **elementary interpretation**:
 - Start with **zero** flow.
 - Repeat: choose an appropriate path from **s** to **t** , and **increase** flow along the edges of this path **as much as** possible.

A Flow Example


(a)




(b)



(c)



(d)





A Closer Look at the Algorithm

- There is just one complication.

What if we choose a path that blocks all other paths?

- Simplex gets around this problem by also allowing paths to cancel existing flow.



A Closer Look at the Algorithm

- To summarize, in each iteration simplex looks for an $s - t$ path whose edges (u, v) can be of two types:
 - ① (u, v) is in the original network, and is not yet at **full capacity**.
 - ② The **reverse** edge (v, u) is in the original network, and there is **some flow** along it.
- If the current flow is f , then in the first case, edge (u, v) can handle up to $c_{uv} - f_{uv}$ **additional units** of flow;
- in the second case, up to f_{vu} **additional units** (canceling **all** or **part** of the existing flow on (v, u)).


A Closer Look at the Algorithm

- These flow-increasing opportunities can be captured in a **residual network** $G^f = (V, E^f)$, which has exactly the two types of edges listed, with residual capacities c^f :

$$\begin{cases} c_{uv} - f_{uv} & \text{if } (u, v) \in E \text{ and } f_{uv} < c_{uv} \\ f_{vu} & \text{if } (v, u) \in E \text{ and } f_{vu} > 0 \end{cases}$$



- Thus we can equivalently think of **simplex** as choosing an $s - t$ path in the **residual network**.
- By simulating the behavior of simplex, we get a **direct algorithm for solving max-flow**.
- It proceeds in **iterations**, each time **explicitly constructing** G^f , finding a suitable $s - t$ path in G^f by using, say, a linear-time **breadth-first search**, and **halting** if there is no longer any such path along which flow can be increased.

The Example





A Flow Example





(a)



(b)




A Flow Example



Cuts

A truly remarkable fact:

Not only does simplex **correctly compute** a maximum flow, but it also generates a **short proof of the optimality** of this flow!



An (s, t) -cut partitions the vertices into two disjoint groups L and R such that $s \in L$ and $t \in R$. Its capacity is the total capacity of the edges from L to R , and as argued previously, is an **upper bound** on any flow: Pick any flow f and any (s, t) -cut (L, R) . Then $\text{size}(f) \leq \text{capacity}(L, R)$.

A Certificate of Optimality

Theorem (Max-flow min-cut)

The size of the maximum flow in a network equals the capacity of the smallest (s, t) -cut.

A Certificate of Optimality

Proof

- Suppose f is the final flow when the algorithm terminates.
- We know that node t is no longer reachable from s in the residual network G^f .
- Let L be the nodes that are reachable from s in G^f , and let $R = V \setminus L$ be the rest of the nodes.
- We claim that $\text{size}(f) = \text{capacity}(L, R)$.
- To see this, observe that by the way L is defined, any edge going from L to R must be at **full capacity** (in the current flow f), and any edge from R to L must have **zero flow**.
- Therefore the net flow across (L, R) is exactly the capacity of the cut.

Efficiency

- Each iteration is efficient, requiring $O(|E|)$ time if a **DFS** or **BFS** is used to find an $s - t$ path.

But how many iterations are there?

- Suppose all edges in the original network have **integer capacities** $\leq C$. Then on each **iteration** of the algorithm, the flow is always an integer and increases by an **integer amount**.
- Therefore, since the maximum flow is at most $C|E|$, the number of iterations is at most this much.
- If paths are chosen in a sensible manner - in particular, by using a **BFS**, which finds the path with the **fewest edges** - then the number of iterations is at most $O(|V| \cdot |E|)$, no matter what the capacities are. *Edmonds-Karp algorithm*
- This latter bound gives an overall running time of $O(|V| \cdot |E|^2)$ for maximum flow.

Efficiency

Lemma:

If the *Edmonds-Karp algorithm* is run on a flow network $G = (V, E)$ with source s and sink t , then for all vertices $v \in V - \{s, t\}$, the shortest path distance $\delta_f(s, v)$ in the residual network G_f increases monotonically with each flow augmentation.

proof:

- Prove by contradiction. Let f be the flow just before the first augmentation that decreases some shortest path distance, and let f' be the flow just afterward.
- Let v be the shortest vertex with the minimum $\delta_{f'}(s, v)$ whose distance was decreased by the augmentation, so that $\delta_{f'}(s, v) < \delta_f(s, v)$. Let $s \rightsquigarrow u \rightarrow v$ be a shortest path from s to v in $G_{f'}$, so that $(u, v) \in E_{f'}$ and

$$\delta_{f'}(s, u) = \delta_{f'}(s, v) - 1$$

Efficiency

proof:

- Because of how we chose v , we know that the distance label of vertex u did not decrease, i.e.,

$$\delta_{f'}(s, u) \geq \delta_f(s, u)$$

- $(u, v) \notin E_f$.
 - otherwise $\delta_f(s, v) \leq \delta_f(s, u) + 1 \leq \delta_{f'}(s, u) + 1 = \delta_{f'}(s, v)$
- $(u, v) \notin E_f$ but $(u, v) \in E_{f'}$? The augmentation must have increased the flow from v to u .
- The Edmonds-Karp algorithm always augments flow along shortest paths, and therefore the shortest path from s to u in G_f has (v, u) as its last edge.
- $\delta_f(s, v) = \delta_f(s, u) - 1 \leq \delta_{f'}(s, u) - 1 = \delta_{f'}(s, v) - 2$.
- Contradict to $\delta_{f'}(s, v) < \delta_f(s, v)$.

Efficiency

Theorem:

If the *Edmonds-Karp algorithm* is run on a flow network $G = (V, E)$ with source s and sink t , then the total number of flow augmentations performed by the algorithm is $O(V \cdot E)$.

Efficiency

proof:

An edge (u, v) in a residual network G_f is **critical** p if the residual capacity of p is the residual capacity of (u, v) , $c_f(p) = c_f(u, v)$. **Each of the $|E|$ edges can become critical at most $|V|/2 - 1$ times.** Let $(u, v) \in E$,

- Since augmenting paths are shortest paths, when (u, v) is critical for the first time, we have

$$\delta_f(s, v) = \delta_f(s, u) + 1$$

- It reappears after the flow from u to v is decreased, which occurs only if (v, u) appears on an augmenting path. We have

$$\delta_{f'}(s, u) = \delta_{f'}(s, v) + 1$$

Efficiency

proof:

- $\delta_{f'}(s, u) \geq \delta_f(s, u) + 2$.
- The distance of u from the source is initially at least 0. Until u becomes unreachable from the source, if ever, its distance is at most $|V| - 2$. Thus, (u, v) can become critical at most $(|V| - 2)/2 = |V|/2 - 1$ times.
- Since there are $O(E)$ pairs of vertices that can have an edge between them in a residual graph, the total number of critical edges during the entire execution of the Edmonds-Karp algorithm is $O(V \cdot E)$.

Min-Max Relations in LP

LP for Max Flow

$$\max f_{ts}$$

$$f_{ij} \leq c_{ij} \quad (i,j) \in E$$

$$\sum_{(w,i) \in E} f_{wi} - \sum_{(i,z) \in E} f_{iz} \leq 0 \quad i \in V$$

$$f_{ij} \geq 0 \quad (i,j) \in E$$

Duality

Primal LP

$$\begin{aligned}\max \quad & c^T \mathbf{x} \\ A\mathbf{x} & \leq b \\ \mathbf{x} & \geq 0\end{aligned}$$

Dual LP

$$\begin{aligned}\min \quad & \mathbf{y}^T b \\ \mathbf{y}^T A^T & \geq c^T \\ \mathbf{y} & \geq 0\end{aligned}$$

Primal LP:

$$\begin{aligned}\max \quad & c_1x_1 + \cdots + c_nx_n \\ a_{i1}x_1 + \cdots + a_{in}x_n & \leq b_i \quad \text{for } i \in I \\ a_{i1}x_1 + \cdots + a_{in}x_n & = b_i \quad \text{for } i \in E \\ x_j & \geq 0 \quad \text{for } j \in N\end{aligned}$$

Dual LP:

$$\begin{aligned}\min \quad & b_1y_1 + \cdots + b_my_m \\ a_{1j}y_1 + \cdots + a_{mj}y_m & \geq c_j \quad \text{for } j \in N \\ a_{1j}y_1 + \cdots + a_{mj}y_m & = c_j \quad \text{for } j \notin N \\ y_i & \geq 0 \quad \text{for } i \in I\end{aligned}$$

LP-Duality

$$\max f_{ts}$$

$$f_{ij} \leq c_{ij} \quad (i,j) \in E$$

$$\sum_{(w,i) \in E} f_{wi} - \sum_{(i,z) \in E} f_{iz} \leq 0 \quad i \in V$$

$$f_{ij} \geq 0 \quad (i,j) \in E$$

$$\min \sum_{(i,j) \in E} c_{ij} d_{ij}$$

$$d_{ij} - p_i + p_j \geq 0 \quad (i,j) \in E$$

$$p_s - p_t \geq 1$$

$$d_{ij} \geq 0 \quad (i,j) \in E$$

$$p_i \geq 0 \quad i \in V$$

Explanation of the Dual

$$\min \sum_{(i,j) \in E} c_{ij} d_{ij}$$

$$d_{ij} - p_i + p_j \geq 0 \quad (i,j) \in E$$

$$p_s - p_t \geq 1$$

$$d_{ij} \in \{0, 1\} \quad (i,j) \in E$$

$$p_i \in \{0, 1\} \quad i \in V$$

- To obtain the dual program we introduce variables d_{ij} and p_i corresponding to the two types of inequalities in the primal.
 - d_{ij} : distance labels on edges;
 - p_i : potentials on nodes.

Integer Program

$$\min \sum_{(i,j) \in E} c_{ij} d_{ij}$$

$$d_{ij} - p_i + p_j \geq 0 \quad (i,j) \in E$$

$$p_s - p_t \geq 1$$

$$d_{ij} \in \{0, 1\} \quad (i,j) \in E$$

$$p_i \in \{0, 1\} \quad i \in V$$

- Let $(\mathbf{d}^*, \mathbf{p}^*)$ be an **optimal solution** to this integer program.
- The only way to satisfy the inequality $p_s^* - p_t^* \geq 1$ with a 0/1 substitution is to set $p_s^* = 1$ and $p_t^* = 0$.
- This solution naturally defines an $s - t$ cut (X, \bar{X}) , where X is the set of potential 1 nodes, and \bar{X} the set of potential 0 nodes.

Integer Program

$$\min \sum_{(i,j) \in E} c_{ij} d_{ij}$$

$$d_{ij} - p_i + p_j \geq 0 \quad (i,j) \in E$$

$$p_s - p_t \geq 1$$

$$d_{ij} \in \{0, 1\} \quad (i,j) \in E$$

$$p_i \in \{0, 1\} \quad i \in V$$

- Consider an edge (i,j) with $i \in X$ and $j \in \bar{X}$, Since $p_i^* = 1$ and $p_j^* = 0$, and thus $d_{ij}^* = 1$.
- The distance label for each of the remaining edges can be set to either 0 or 1 without violating the first constraints; however in order to **minimize** the objective function value it must be set to 0.
- The objective function value is precisely the **capacity** of the cut (X, \bar{X}) , and hence (X, \bar{X}) must be a **minimum** $s - t$ cut.

Relaxation of the Integer Program

- The **integer program** is a formulation of the **minimum $s - t$ cut** problem.
- The **dual program** can be viewed as a **relaxation** of the integer program where the integrality constraint on the variables is dropped.
- This leads to the constraints $1 \geq d_{ij} \geq 0$ for $(i, j) \in E$ and $1 \geq p_i \geq 0$ for $i \in V$.
- The **upper bound constraints** on the variables are redundant; their omission cannot give a better solution.
- Dropping these constraints gives the dual program in the form given above. We will say that this program is the **LP relaxation** of the **integer program**.

Relaxation of the Integer Program

- In principle, the best **fractional** $s - t$ cut could have lower capacity than the best integral cut. Surprisingly enough, this does not happen.
- From **linear programming** theory we know that for any objective function, i.e., assignment of capacities to the edges of G , there is a vertex solution that is **optimal**.
- Now, it can be proven that each vertex solution is integral, with each coordinate being 0 or 1 .
- This follows from the fact that the constraint matrix of this program is totally unimodular, Thus, the dual program always has an **integral optimal solution**.

Referred Materials

- Content of this lecture comes from Section 7.2 in [DPV07], Section 26.1 and 26.2 in [CLRS09] and Section 12.2 in [Vaz04].